

# Golden Metrics

## With Prometheus and Grafana

Johnathan Fercher

# Golden Metrics

With Prometheus and Grafana + Datadog

Johnathan Fercher

## INDEX

### **Introduction**

Observability

Golden Signals

### **Prometheus**

Go API

Agent

### **Grafana**

Dashboards & Troubleshooting

### **Datadog**

Dashboards + Costs

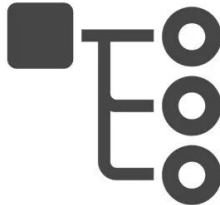
## Introduction

"Observability is the ability to measure the internal states of a system by examining its outputs."

### Three pillars of observability



Metrics



Traces



Logs

Reference:

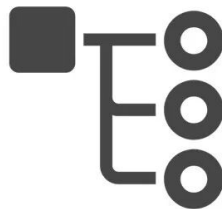
<https://www.confluent.io/learn/observability/#:~:text=What%20is%20Observability%3F,in%20the%20context%20of%20APM.>

**Gympass**

## Introduction

"Observability is the ability to measure the internal states of a system by examining its outputs."

### Three pillars of observability



Reference:

<https://www.confluent.io/learn/observability/#:~:text=What%20is%20Observability%3F,in%20the%20context%20of%20APM.>

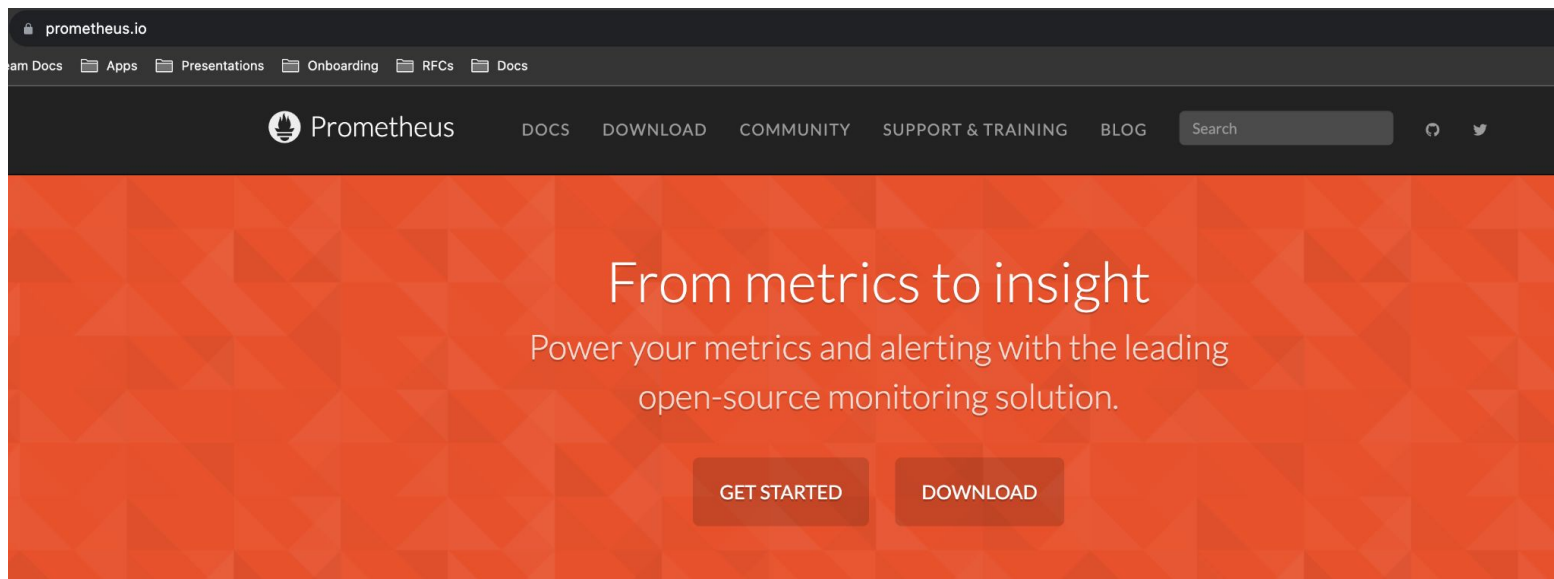
**Gympass**

## Golden Signals

The four golden signals of monitoring are latency, traffic, errors, and saturation. If you can only measure four metrics of your user-facing system, focus on these four.

- Latency
- Traffic
- Errors
- Saturation

# Prometheus



```
func Start() {  
    fmt.Println("starting prometheus")  
    http.Handle("/metrics", promhttp.Handler())  
  
    go func() {  
        http.ListenAndServe(":2112", nil)  
    }()  
    fmt.Println("started prometheus")  
}
```



```
type Metrics struct {  
    // Metric  
    Latency float64  
  
    // Labels  
    Endpoint      string  
    Verb          string  
    Pattern       string  
    ResponseCode  int  
    Failed        bool  
    Error         string  
    HasAvailabilityError bool  
    HasReliabilityError bool  
}
```

## Agent

```
func Send(metrics Metrics) {
    labels := map[string]string{
        endpoint:    metrics.Endpoint,
        verb:        metrics.Verb,
        pattern:     metrics.Pattern,
        responseCode: fmt.Sprintf("%d", metrics.ResponseCode),
        failed:      fmt.Sprintf("%v", metrics.Failed),
        error:       metrics.Error,
        isAvailabilityError: fmt.Sprintf("%v", metrics.HasAvailabilityError),
        isReliabilityError:  fmt.Sprintf("%v", metrics.HasReliabilityError),
    }

    countermetrics.Increment(countermetrics.Metric{
        Name:    endpointRequestCounter,
        Labels:  labels,
    })

    histogrammetrics.Observe(histogrammetrics.Metric{
        Name:    endpointRequestLatency,
        Value:   float64(metrics.Latency),
        Labels:  map[string]string{
            endpoint: metrics.Endpoint,
        },
    },
    })
}
```

Reference: <https://johnfercher.medium.com/go-observabilidade-739b6d6b649c>

## Agent localhost/metrics

```
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 0
go_gc_duration_seconds{quantile="0.25"} 0
go_gc_duration_seconds{quantile="0.5"} 0
go_gc_duration_seconds{quantile="0.75"} 0
go_gc_duration_seconds{quantile="1"} 0
go_gc_duration_seconds_sum 0
go_gc_duration_seconds_count 0
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 9
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
go_info{version="go1.20.2"} 1
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go_memstats_alloc_bytes 1.08112e+06
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
# TYPE go_memstats_alloc_bytes_total counter
go_memstats_alloc_bytes_total 1.08112e+06
# HELP go_memstats_buck_hash_sys_bytes Number of bytes used by the profiling bucket
# TYPE go_memstats_buck_hash_sys_bytes gauge
go_memstats_buck_hash_sys_bytes 1.447357e+06
# HELP go_memstats_frees_total Total number of frees.
# TYPE go_memstats_frees_total counter
go_memstats_frees_total 360
# HELP go_memstats_gc_sys_bytes Number of bytes used for garbage collection system n
# TYPE go_memstats_gc_sys_bytes gauge
go_memstats_gc_sys_bytes 5.707792e+06
```

Reference: <https://johnfercher.medium.com/go-observabilidade-739b6d6b649c>

## Types

- Counter
  - **Accumulative values**
  - **Only increase or reset**
  - Ex: 10 request on endpoint X in the last 5 minutes
  - Ex: 5 errors on endpoint Y in the last 10 minutes
- Gauge
  - **Current values**
  - **Can increase or decrease**
  - Ex: PostgreSQL have 5 connections enabled
  - Ex: Errors threshold is in 10% to circuit-breaker
- Histogram
  - **Sample values in buckets**
  - How many milliseconds does an endpoint spent to return on average?
- Summary
  - **Similar to histogram**

## Your **LGTM** Observability stack

Get there much faster. From dashboards to centralized observability.



### The (actually useful) free forever plan

Grafana, of course +

10K series Prometheus metrics,  
50GB logs, 50GB traces, 50GB profiles,  
500VUh k6 testing

Create free account

(No credit card required)

Your data and tools



Your applications



Your infrastructure



Metrics



Logs



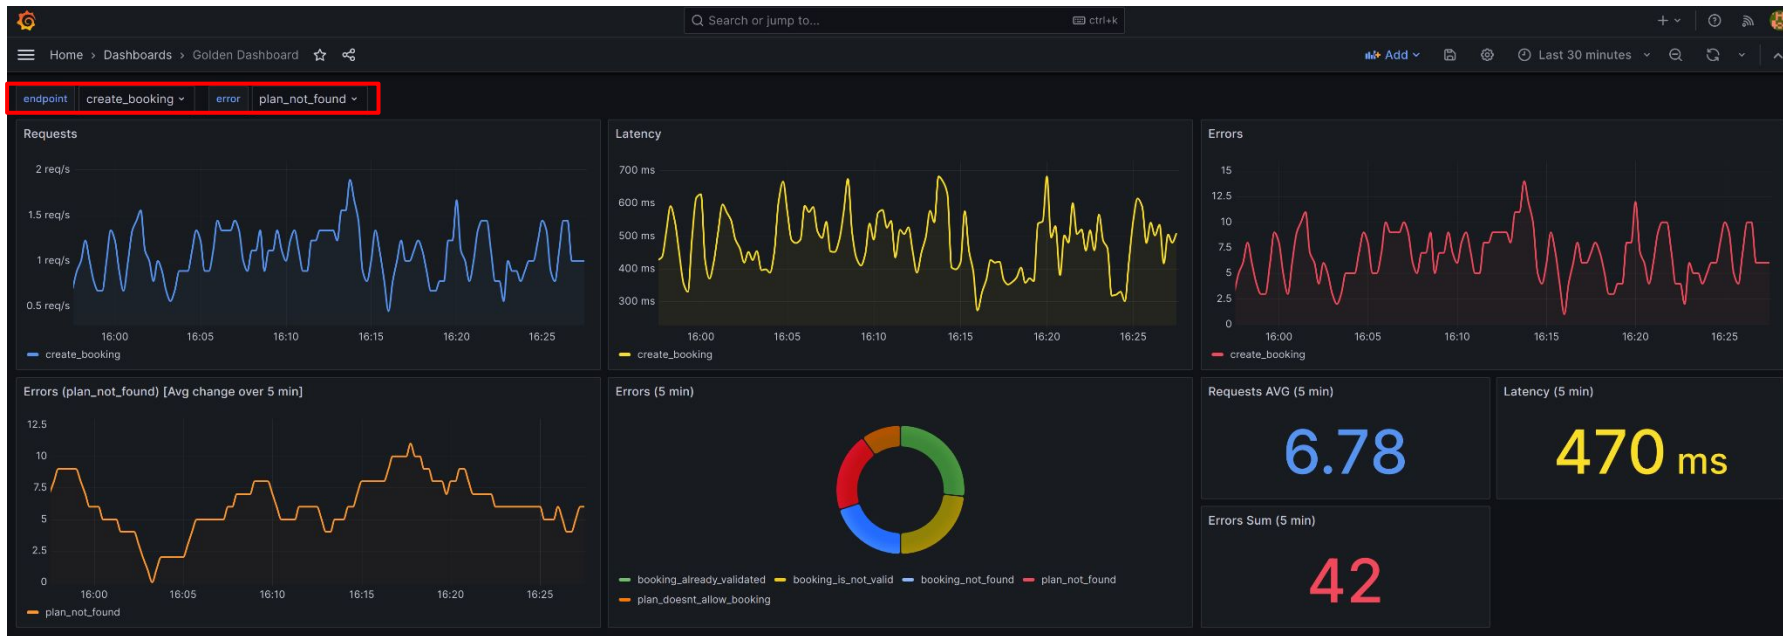
Traces



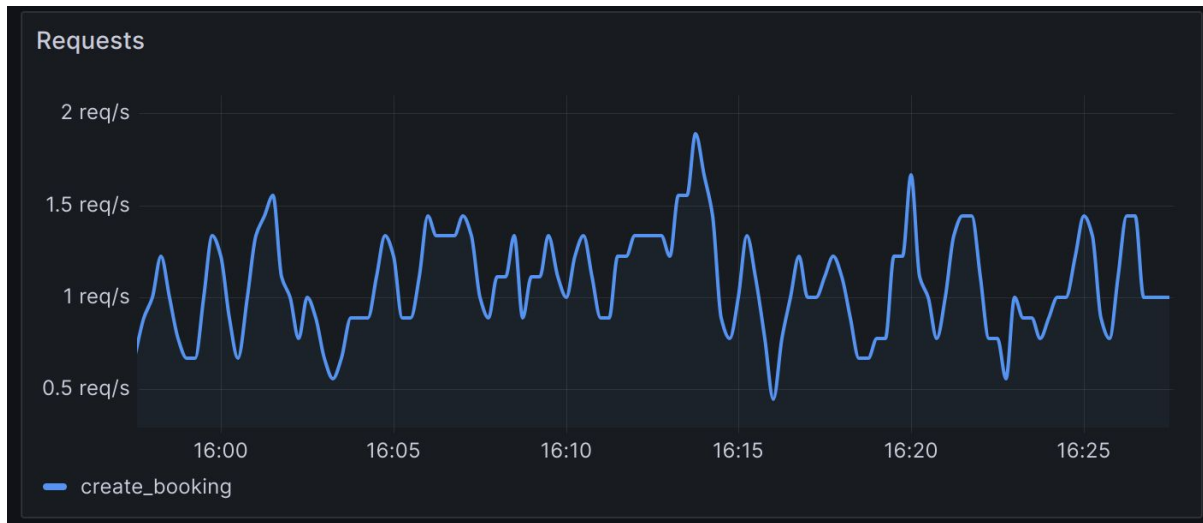
# Dashboard



# Dashboard



## Requests



Requests AVG (5 min)

6.78



## Latency



Latency (5 min)

470 ms

## Errors



Errors Sum (5 min)

42

## Errors



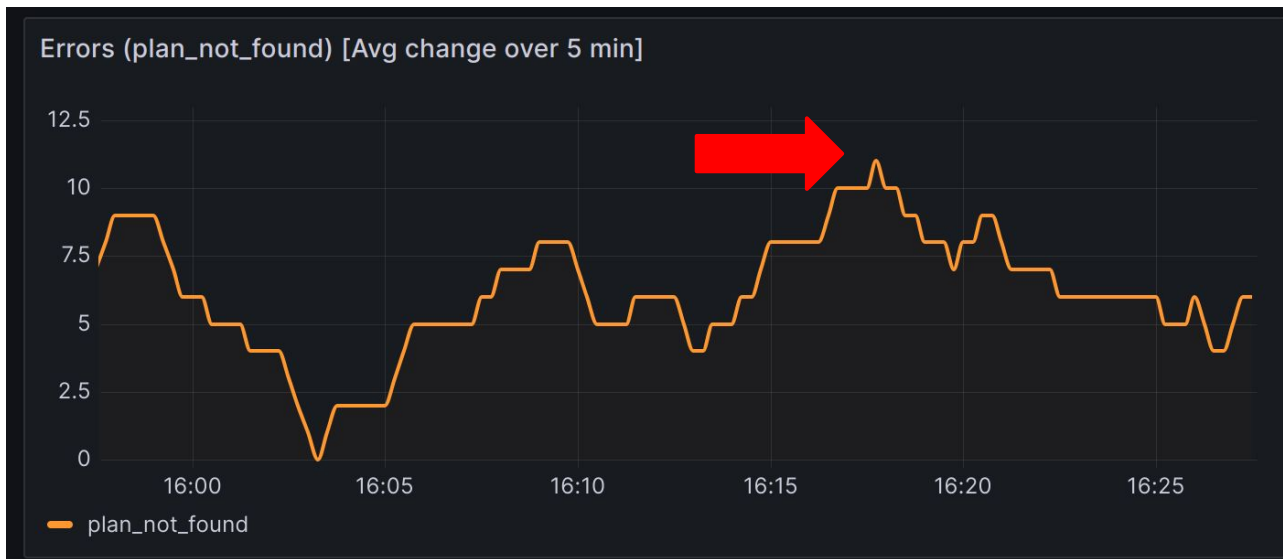
Errors Sum (5 min)

42

## Errors

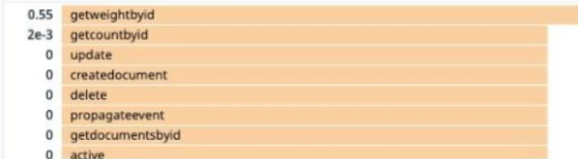


## Errors

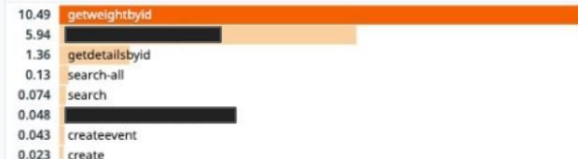


# Datadog

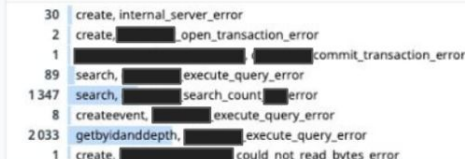
Relative Increase of Errors (%)



Relative Errors (%)



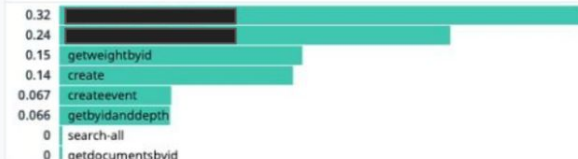
Errors Total



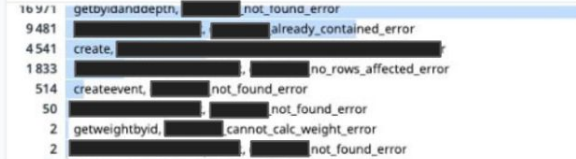
Relative Increase of Attention Signals (%)



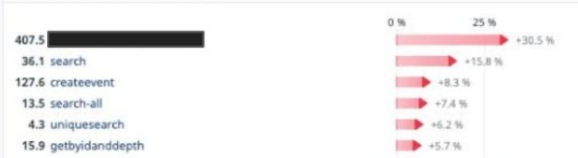
Relative Attention Signals (%)



Attention Signals Total



Relative Increase of Latency



Absolute AVG Latency



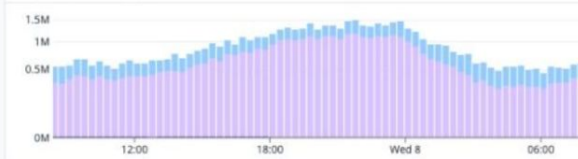
Requests Total



Relative Increase of Requests



Requests Per Scope

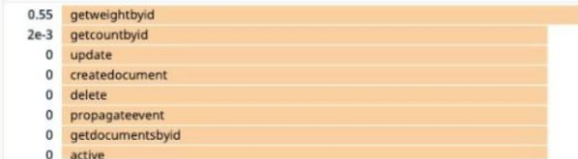


Total Requests

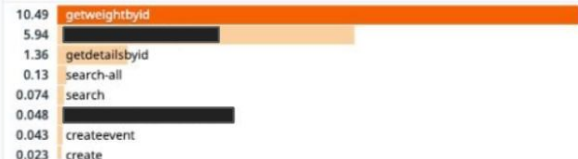
63.17M

# Datadog

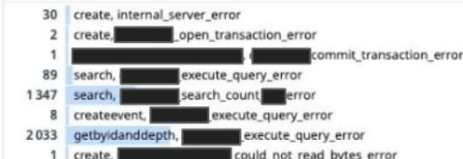
Relative Increase of Errors (%)



Relative Errors (%)



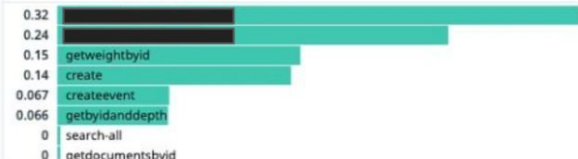
Errors Total



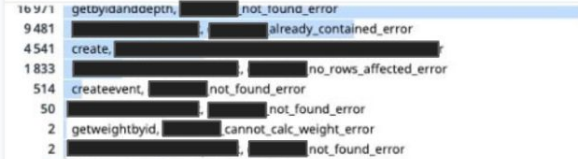
Relative Increase of Attention Signals (%)



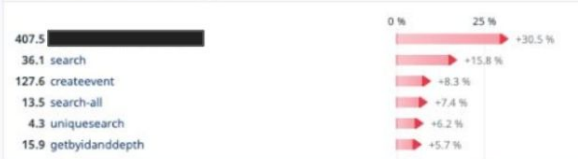
Relative Attention Signals (%)



Attention Signals Total



Relative Increase of Latency



Absolute AVG Latency



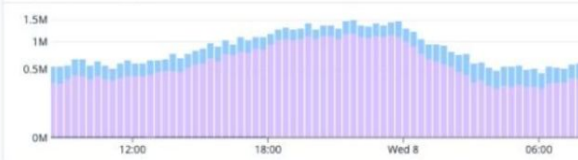
Requests Total



Relative Increase of Requests



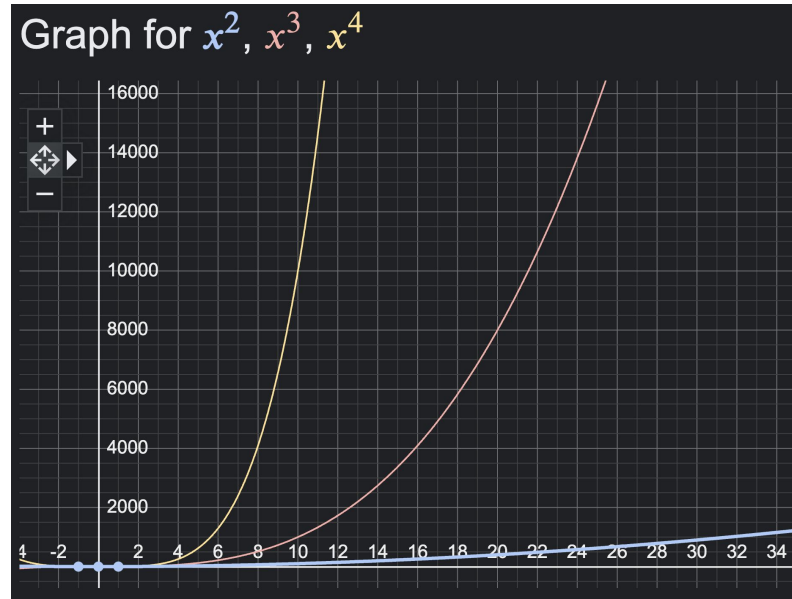
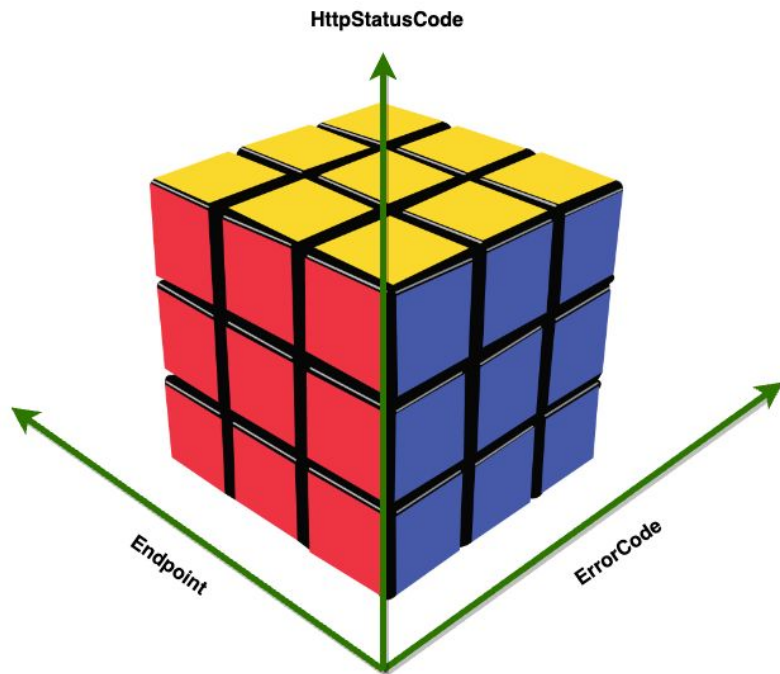
Requests Per Scope



Total Requests

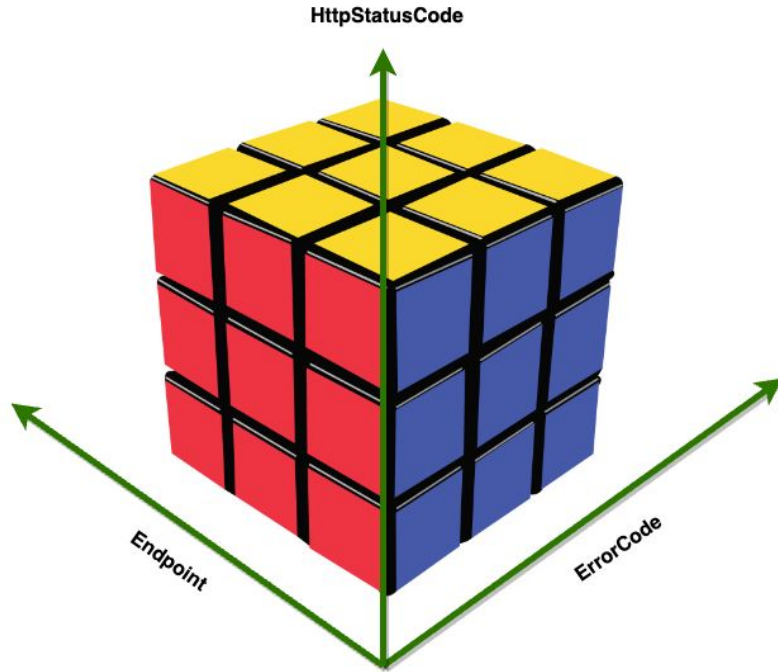


## Is This Expensive?



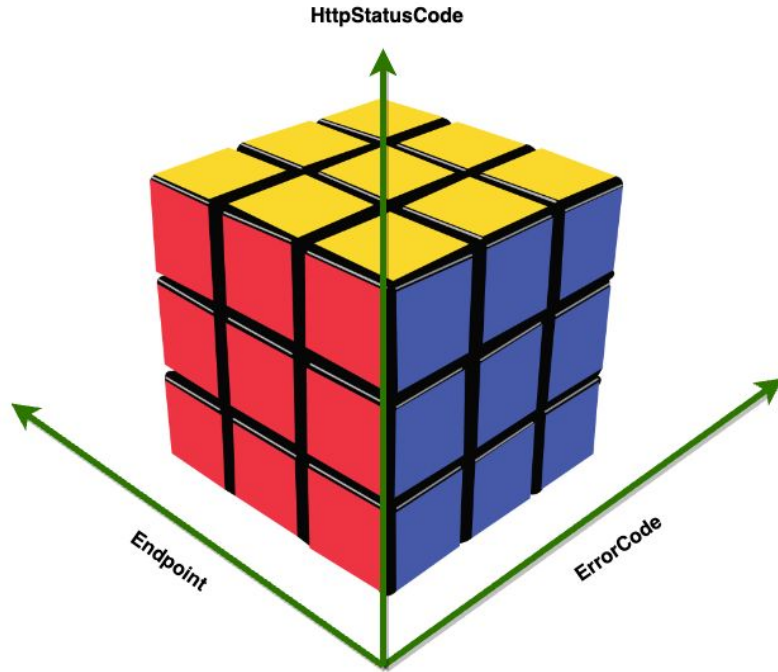


**Is This Expensive?**



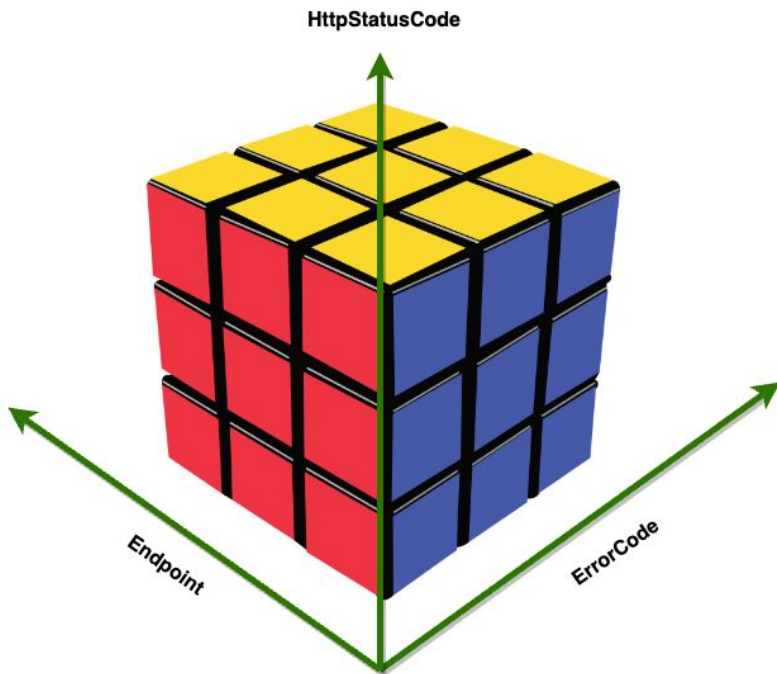
**Cost = Space + Write + Read**

Is This Expensive?



Cost = **Space** + Write + Read

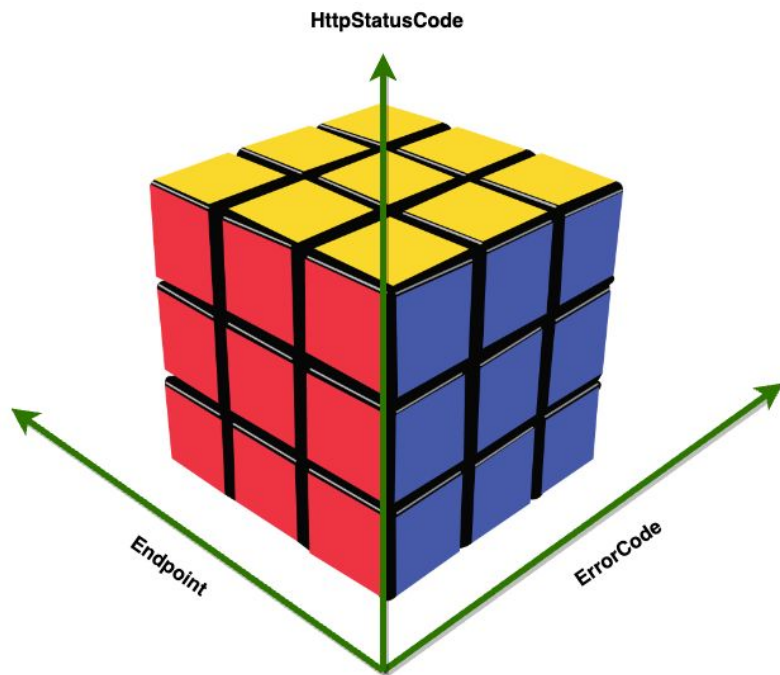
## Is This Expensive?



Cost = **Space** + Write + Read

Space =  $\text{Len}(\text{Endpoint}) * \text{Len}(\text{ErrorCode}) * \text{Len}(\text{HttpStatusCode}) * \dots$

## Is This Expensive?



Cost = **Space** + Write + Read

Space =  $\text{Len}(\text{Endpoint}) * \text{Len}(\text{ErrorCode}) * \text{Len}(\text{HttpStatusCode}) * \dots$

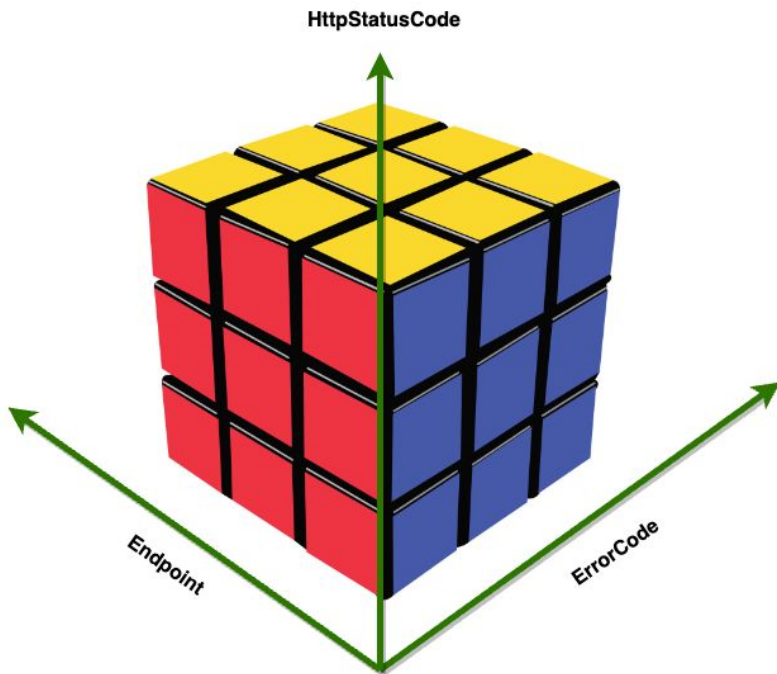
### Small API

Space = 12 endpoints \* 25 error codes \* 10 http status

Space = 3.000

(This would not increase by time)

## Is This Expensive?



Cost = **Space** + Write + Read

Space =  $\text{Len}(\text{Endpoint}) * \text{Len}(\text{ErrorCode}) * \text{Len}(\text{HttpStatusCode}) * \dots$

### Small API

Space = 12 endpoints \* 25 error codes \* 10 http status

Space = 3.000

(This would not increase by time)

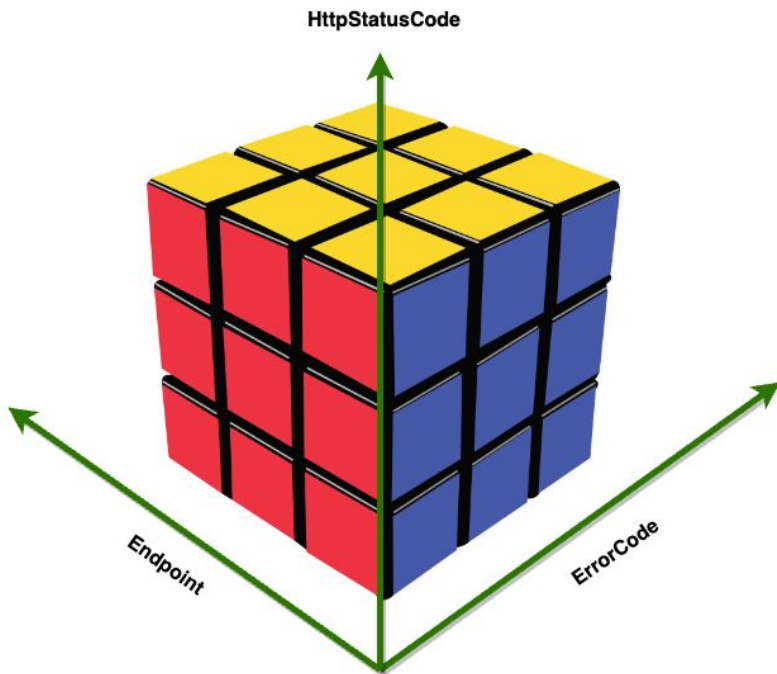
### Big API

Space = 50 endpoints \* 100 error codes \* 10 http status

Space = 50.000

(This would not increase by time)

## Is This Expensive?



Cost = **Space** + Write + Read

Space =  $\text{Len}(\text{Endpoint}) * \text{Len}(\text{ErrorCode}) * \text{Len}(\text{HttpStatusCode}) * \dots$

### Small API

Space = 12 endpoints \* 25 error codes \* 10 http status

Space = 3.000

(This would not increase by time)

### Big API

Space = 50 endpoints \* 100 error codes \* 10 http status

Space = 50.000

(This would not increase by time)

### Wrong Usage (Small API)

Space = 12 endpoints \* **25000 booking\_id** \* 10 http status

Space = 3.000.000

(This would increase **daily** by time)

Gympass

**Thank You**